

Crop Yield Intelligence System (AgriPredict)

A Machine Learning-Based Approach to Agricultural Yield Prediction

Kush

Department of Computer Science & Engineering
Lovely Professional University
Phagwara, India
sonisingh0107@gmail.com

Nishar

Department of Computer Science and Engineering
Lovely Professional University
Phagwara, India
sonisingh0107@gmail.com

April 2026

Abstract

Forecasting agricultural output has become central to tackling the growing challenge of global food security. With the world's population on a steady rise, there is an ever-increasing need for farming practices grounded in data and supported by modern analytical tools. This paper introduces AgriPredict, a Crop Yield Intelligence System built upon supervised machine learning to estimate crop productivity in tons per hectare using a broad range of environmental and agronomic inputs.

The system incorporates three predictive models — Random Forest, XGBoost (Extreme Gradient Boosting), and a hyperparameter-optimized RandomForest variant — all trained on a large-scale synthetic dataset comprising one million agricultural records that cover four geographic regions, six soil categories, six crop varieties, and three weather scenarios.

Among the models evaluated, XGBoost delivered the strongest results, recording a coefficient of determination (R^2) of 0.9117 and a mean absolute error (MAE) of 0.402 tons per hectare. These figures indicate that the model accounts for more than 91% of the variation observed in crop yield outcomes. An analysis of feature contributions shows that precipitation is the dominant factor at 60.7%, with fertilizer application contributing 20.1% and irrigation accounting for 12.6%.

The platform offers a browser-based interface through which users can enter field-specific conditions and instantly obtain yield forecasts, making it practical for farmers, planners, and agricultural researchers. This paper covers the system design, data preparation workflow, rationale for model selection, the training and evaluation approach, and the broader implications of real-world deployment.

I. INTRODUCTION

Agriculture remains the backbone of many developing economies. In India, for instance, more than 58% of rural households depend on farming for their livelihood. Despite its importance, the sector faces constant challenges: unpredictable weather, uneven soil conditions, and limited access to structured planning tools. For smallholder farmers, the absence of reliable yield forecasting often leads to poor resource management, overuse of fertilizers and water, financial uncertainty, and threats to food availability.

Traditional yield estimation methods — manual inspections, expert judgment, or rule-of-thumb models — are slow and prone to error. Machine learning offers a transformative alternative. These algorithms can process vast amounts of data and uncover complex, non-linear relationships between farming conditions and harvest outcomes that traditional statistical methods often miss.

AgriPredict was developed to forecast yields using nine input parameters: rainfall, temperature, soil type, region, crop variety, weather, fertilizer use, irrigation, and days to harvest. Built on a dataset of one million records, the system provides farmers and policymakers with a real-time planning tool that reduces the cost and time of manual estimation. On a larger scale, it supports the UN's Sustainable Development Goal 2 (Zero Hunger) by promoting smarter, data-driven farming decisions.

II. LITERATURE REVIEW

Over the past decade, machine learning has become central to agricultural yield forecasting. Advances in remote sensing, IoT sensors, and digital farming records have fueled this growth. Researchers have tested regression models, decision trees, ensemble methods, and deep learning, each with strengths depending on the dataset.

Liakos et al. (2018) reviewed machine learning in precision farming and found that ensemble methods like Random Forest often outperform single models by reducing overfitting. Jeong et al. (2016) demonstrated Random Forest's effectiveness in predicting soybean and maize yields using satellite imagery.

XGBoost, introduced by Chen and Guestrin (2016), has become a leading algorithm for structured prediction tasks. Its gradient boosting framework with built-in regularization often achieves better generalization than Random Forest. Filippi et al. (2019) applied XGBoost to wheat yield prediction and achieved R^2 values above 0.88, similar to the results in this study.

Deep learning models such as LSTMs have also been explored for sequential yield forecasting. While powerful, they require larger datasets and more computing resources. For structured datasets like the one used here, tree-based methods remain the most practical choice. Many existing systems are limited to specific crops or regions, but

AgriPredict aims to be broader, covering multiple crops and environments.

III. SYSTEM ARCHITECTURE

AgriPredict is developed as a fully-integrated intelligent solution combining data ingestion and preprocessing pipeline, backend model training and evaluation, and front-end web-based user interface for communication with users. The three-tier architecture is applied to separate responsibilities, make maintenance easier, and facilitate extensibility of the system in future.

A. Data Layer

Data layer contains a structured CSV dataset which comprises one million agriculture-related samples. For each sample there is information about environmental and agronomical factors as well as the obtained agricultural output measured in tons per hectare. Data layer is responsible for all data processing before passing to model layer. It involves reading dataset from the file, handling missing data, transforming categorical factors into numerals, normalizing numerical values, and splitting the data into train and test subsets.

B. Model Layer

The core of AgriPredict is composed of three regression models: Baseline Random Forest model, XGBoost Regressor, and optimized Random Forest model. Both model types are implemented using scikit-learn and XGBoost libraries in Python. Upon being trained, each model is serialized to a file using joblib library. These metrics are measured on the unseen test partition and documented for performance evaluation.

C. Presentation Layer

The presentation layer consists of a React.js web application served locally from a development server listening at port 5000. The application comprises five tabs: Predict Yield, Analysis, Model Details, Retrain Model, and About Project. When making predictions, users provide values for all nine inputs and make an inference call, receiving their results in return. The Analysis tab displays interactive bar plots created by Recharts library to illustrate yield differences based on soil type, location, and application of fertilizers or irrigation systems. Model Details tab provides a tool for comparative analysis of model-based metrics and learning curves.

D. Input–Processing–Output Pipeline

The whole process starts when users submit farm-level data via user interfaces. After validation, encoding, and normalization of inputs using the same preprocessing pipeline used in the training step, the processed feature vector is provided to the selected model, yielding an estimated value for the output variable. Users receive their predictions, which are displayed on the screen and marked as points on the Actual vs. Predicted scatter plot. The constant badge of the model in the navigation header includes the name of the current model, its R^2 value, and the size of the entire training data.

IV. MACHINE LEARNING MODELS

A. Random Forest Regressor

Random Forest is a decision tree ensemble algorithm based on the bootstrap aggregation technique. In this case, a large number of trees are generated during training, and their output is aggregated by calculating the average value of predictions. The key element of this technique lies in building a tree using a subset of training data samples. Additionally, the Random Forest model uses the randomness property when choosing feature splits, thereby preventing overfitting. The accuracy of the Random Forest model within AgriPredict was 0.8989, and the root mean square error was 0.5386. Cross-validation results for the Random Forest model showed a mean R^2 value of 0.8990. This indicated consistent generalization behavior. In addition, the learning curves of the Random Forest model showed stable results at various training data sizes. At the same time, the training and validation accuracies were almost identical, stabilizing at 0.90, indicating the regularity of the model's performance.

B. XGBoost Regressor

XGBoost stands for "Extreme Gradient Boosting," which is a sequential ensemble technique whereby trees are trained sequentially and each subsequent tree tries to learn from the mistakes made by previous trees. Unlike bagging methods, boosting algorithms allocate more weight to those data instances that have been poorly estimated and thus continuously concentrate on the hardest cases. Moreover, in addition to using L1 and L2 regularization right in the loss function to reduce overfitting, XGBoost is also capable of working with sparse data. For AgriPredict, XGBoost was tuned up with 100 estimators and max_depth of 20. After training on 800,000 rows and validation on another 200,000, XGBoost demonstrated the best result according to all quality measures: $R^2 = 0.9117$ and MAE = 0.402 tons per hectare. Therefore, XGBoost served as an inference algorithm, and it is illustrated as such in the navigation bar.

C. Tuned Random Forest

Since structured hyperparameter optimization would help find out whether Random Forest could bridge the performance disparity compared to XGBoost, an optimized version of the Random Forest algorithm was developed as well. The hyperparameter optimization was done through cross-validation, where a grid search included several key hyperparameters like the number of trees, maximum tree depth, minimal sample number at split level, and the number of randomly selected features at each tree node. Although the tuned Random Forest model provided a marginal improvement compared to its untuned counterpart, it still lagged behind the performance of XGBoost, which is designed on the basis of sequentially constructed models that are more expressive than Random Forest..

V. WORKING OF THE MODEL


A. Preprocessing Pipeline

Prior to entering any model, the initial raw agricultural data goes through a well-defined series of multiple preprocessing steps. The categorical attributes, namely Crop Type, Region, Soil Type, and Weather Condition, are

transformed to their numerical equivalents via label encoding, whereby each category is assigned an integer value corresponding to its unique label. The numeric attributes such as Rainfall in mm, Temperature in °C, Days until Harvest, Fertilizer used (a binary variable), and Irrigation used (another binary variable) are normalized via min-max scaling to fit between zero and one. This is done to mitigate the impact that attributes with large numeric ranges may have on the learning process of the model..

B. Feature Engineering and Selection

Feature contribution values were obtained using the in-built feature attribution approach of XGBoost to determine which features were the driving forces behind the predictions. This yielded the following results. First, it was determined that Rainfall is the most important predictor, making up 60.7% of the weight of the predictors, thereby proving beyond doubt that rainfall is the main contributor to crop productivity in this setting. Second, Fertilizer comes in at second position with 20.1% of the predictor weight, and this is attributed to the agricultural science understanding that nutrient provision greatly enhances crop production. Third, Irrigation contributes 12.6% towards prediction weight, confirming the significance of irrigation in compensating for irregular rainfalls. Finally, Temperature and Days to Harvest make contributions of 2.8% and 1.6% respectively, and the remaining 2.3% is made up by the categorical predictors.



C. Prediction Process

Once the prediction query is initiated, the input parameters via the web interface undergo processing through the preprocessor pipeline. These features will be encoded using the same encodings and normalized based on the same normalization parameters used during the training stage. This ensures complete parity between the transformations done during the training phase and those applied to the features before feeding them into the machine learning model. This feature vector will then be sent to the XGBoost model trained on our dataset and processed by the 100 gradient-boosted decision trees. Ultimately, the model will return a single continuous numeric value indicating the yield in tons per hectare..

VI. DATASET AND PARAMETERS

A. Dataset Overview

The database used for training and testing the AgriPredict model includes one million synthetic records in agriculture. Following the elimination of redundant information and those with improbable values during the data cleaning phase, 999,769 observations were left for subsequent processing. This database contains four geographical regions (North, East, South, and West), six soil types (Chalky, Clay, Loam, Peaty, Sandy, and Silt), six crops (Soybean, Wheat, Rice, Cotton, Maize, and Barley), and three weather conditions (Sunny, Rainy, and Cloudy). Yield values range between 0 and 9.96 tons per hectare, with an average of 4.65 tons per hectare for the entire database.

B. Input Features

There are nine input features that make up one record in the data set. These include: (i) The type of crop, which refers to the specific variety grown and is a categorical input feature; (ii) The region, representing the geographical division within which the farm is located, which indirectly represents climate and terrain features; (iii) Soil type, representing the texture of the soil in the farmland, thus impacting its water retention capabilities; (iv) Weather condition, representing the predominant weather condition in the particular season (Sunny, Rainy, or Cloudy); (v) Rainfall in millimeters, varying from 100 mm in dry seasons to 1000 mm in wet seasons; (vi) Temperature in degree centigrade, which varies between 15°C and 40°C; (vii) Days to harvest, measuring the total growing period ranging between 60 and 149 days; (viii) Application of chemical fertilizers, Yes or No; (ix) Irrigation, Yes or No..

C. Hyperparameters

XGBoost Regressor was initialized with the following parameters: `n_estimators=100` (number of boosting rounds), `max_depth=20` (depth restriction for each tree), `learning_rate=0.1`, `subsample=0.8`, `colsample_bytree=0.8`, `reg_alpha=0.1` (L1 regularization), `reg_lambda=1.0` (L2 regularization). The baseline random forest model was implemented with scikit-learn's default parameters except for `n_estimators` which is equal to 100. Tuned Random Forest was tested under different combinations of parameters including `n_estimators` in [50, 100, 200], `max_depth` in [10, 20, None], and `min_samples_split` in [2, 5, 10]..


VII. DATA SPLITTING AND VALIDATION

The total number of records in the dataset after cleaning was 999,769, which we used to create train and test partitions, following a common convention of using an 80-20 split ratio. This means that we have about 799,815 observations for training and 199,954 observations for testing. To keep our partition process robust, we used random stratification based on overall yield values in both training and test datasets.

While creating train and test partitions was sufficient to evaluate generalization performance for the two models, we further evaluated the performance of these models by applying a technique called five-fold cross-validation. Here, we partition our training set into five equal parts. For each fold in our validation set, we use four other folds to train the model and validate it on the left-out fold. We repeat the

process five times and average our performance scores to obtain a lower variance estimate. The average R^2 obtained through five-fold cross-validation for the XGBoost model was 0.9117, which was very close to the actual test R^2 of 0.9135, implying no overfitting. Similarly, the Random Forest model had an R^2 of 0.8990 for cross-validation..

Learning curve of Random Forest algorithm available under the Model Details tab in the app shows the relationship between R^2 performance and the incremental ratio of training dataset size (from 0% to 9% of total samples). In this case, the score achieved during training, which stayed close to 0.98, and the score obtained during validation, which stabilized at 0.90, did not change throughout all training ratios investigated. Such performance behavior is characteristic of regularization algorithms when learning on large training datasets..



VIII. TRAINING AND TESTING

Training of all three models was done using Python version 3.10, with scikit-learn 1.2 and XGBoost 1.7 serving as the main machine learning libraries. Training of all models took place on a CPU with no GPU acceleration, where XGBoost's inherent ability to utilize multiple threads for constructing trees made training on the one-million row dataset efficient.

The training of the model was done in an organized manner: the CSV file was loaded into a DataFrame, columns of categorical data were encoded using label encoding, numerical data was standardized through min-max scaling, and the features and targets were separated. The models were then trained through a call to their `fit()` function on the training data set. Following the training, the trained model instances were dumped onto the hard drive through `joblib.dump()`.

The evaluation of the model performance was carried out on the hold-out test subset using four metrics commonly used in regression analysis. These were the coefficient of determination (R^2), root mean squared error (RMSE), mean absolute error (MAE), and cross-validation mean R^2 . R^2 indicates the fraction of the variation in the dependent variable accounted for by the predictive power of the model, where a closer proximity to 1.0 signifies a better fit. RMSE has a bias toward larger deviations from predictions because


of the squaring operation, whereas MAE presents a simple metric of prediction deviation in tons per hectare.

Both XGBoost and Random Forest showed better performance than their competitor counterparts in terms of all presented measures. The R^2 achieved by XGBoost was 0.9117 with MAE equaling 0.402 t/ha. On the other hand, the Untuned Random Forest achieved $R^2 = 0.8989$, RMSE = 0.5386, and MAE = 0.4297 t/ha. Thus, the rather close performance of these algorithms allows one to say that the data set is sufficiently structured for the application of machine learning techniques and the upper threshold of the R^2 for tree-based regressions lies somewhere around the interval [0.91, 0.92].

IX. GRAPHS AND RESULTS ANALYSIS

A. Actual vs. Predicted Scatter Plot


The main chart on the Predict Yield sheet is a scatter plot depicting the values of the observed yield versus the values of the predicted one calculated using the XGBoost algorithm on a randomly selected sample of test instances. The horizontal axis indicates the value of the real yield in tons per hectare, while the vertical axis shows the value of the predicted yield. Every dot in the chart depicts one instance of testing. An ideal scenario is when all dots lie on a line that coincides with the 45-degree diagonal. For the AgriPredict company, the majority of the dots in the scatter plot lie near this line, especially in cases where the values of the yields are between 3 and 7 t/ha. There are several deviations of dots below 1 t/ha and above 8 t/ha, which is normal for agricultural modeling. The latest prediction made by the user is represented by a red dot in the chart.



B. Fertilizer and Irrigation Impact

The graph shown below the Analysis tab depicts the average yields of crops in four categories: fields without any fertilizer application, fields using fertilizer applications, fields without any irrigation, and fields with irrigation. The variance between these fields is stark: crops from fields utilizing fertilizer gave an average yield of roughly 5.37 t/ha, while those without had a yield of roughly 3.89 t/ha, a difference of around 38%. Those that were irrigated gave a

yield of roughly 5.28 t/ha, while those not irrigated yielded around 4.03 t/ha, an increase of 31%. These figures provide evidence for the widely known concept in agriculture that nutrient and water management are some of the most crucial determinants in crop yield. The extreme variance seen in fertilization and no fertilization is similar to the variable importance result showing a 20.1% contribution of fertilizers to the model prediction.




C. Yield by Soil Type

The following graph shows the mean yield of each of the six types of soils: Chalky, Clay, Loam, Peaty, Sandy, and Silt. The data lie in a close range, ranging between about 4.50 to 4.66 t/ha. The highest mean yield is recorded in the case of Chalky soil, which has a mean yield of about 4.65 t/ha, while the other two soils, Clay and Silt, have slightly lower values of about 4.63 t/ha. This means that the range being tight means that classification into soil types alone does not play an important role in determining the yield. Rather, soil type acts in combination with other factors such as rain, availability of fertilizers, and irrigation to create the impact that it does.

D. Average Yield by Region

In the comparative table for regions, the East region tops the list for its maximum average yield of about 4.66 tons per hectare, while West comes second with an average yield of 4.65 tons per hectare, followed by South with 4.63 tons per hectare, and North, which yields the least amount of 4.61 tons per hectare. Although these figures may seem insignificant, they do matter significantly when we consider them from the perspective of agriculture as a broad scale strategy. It may be due to the favorable climatic conditions represented in the data.



E. Model Performance Summary

In the Model Details section, there is an elaborate comparison of all the three trained models. From the difference in their R^2 values, one can see how much better XGBoost performs than Random Forest, with an R^2 of 0.9117 against an R^2 of 0.8989 for Random Forest. While the difference may not be noticeable to some, at about 1.3%, this will actually translate to hundreds of thousands more accurate predictions of yields when implemented on a nationwide level. One can tell that Random Forest does not overfit based on its learning curve graph.

X. ADVANTAGES AND LIMITATIONS

A. Advantages

The AgriPredict system offers a variety of important benefits compared to traditional solutions used in the domain of crop yield prediction. First and foremost, this solution is capable of providing real-time prediction, which cannot be done using manual analysis techniques requiring several weeks of time. In addition, the machine learning model in question has been tested on one million examples and is therefore much more generalizable than algorithms designed specifically for a certain crop or region. The fact that the algorithm provides information about feature contributions also proves to be very valuable, as it makes it possible for stakeholders (farmers, policymakers) to comprehend not only what the predicted yield is, but what determines this particular value. The use of the web-based UI eliminates any technical barriers – even a farmer without any background knowledge will be able to plug in some input data and obtain results. Finally, the existence of multiple models on the same platform enables comparisons between different models.

Socially speaking, the AgriPredict software has the potential to play a valuable role in facilitating the efforts of government agencies that work towards ensuring food security by offering an accurate forecast at the national level regarding the production of crops. In addition, financial institutions that provide services such as crop loans to farmers will benefit from the AgriPredict software because it enables them to determine the level of risks involved. NGOs dealing with farming areas that produce little can take advantage of this technology to determine the input that

limits production in these areas. This has something to do with SDG 2 (Zero Hunger)..

B. Limitations

Despite its excellent performance on benchmarks, AgriPredict has several notable limitations. First, the data used in training the model is artificial and not sourced from actual agricultural practices or sensors in the field. While artificial data provides for extensive, controlled testing, it is highly unlikely to capture all the intricacies involved in real-life farming, such as diseases affecting crops, flooding incidents, gradual deterioration of soil quality, and cumulative impacts of climate change. Second, there are just nine predictor variables in the current model, while actual precision agriculture operations consider a variety of additional metrics like soil acidity, organic matter composition, insect threat indicators, and satellite-based metrics. Finally, the system lacks confidence intervals or measures of prediction uncertainty, which can prove extremely useful in risk-sensitive planning scenarios.

XI. FUTURE SCOPE

Several interesting advancements are proposed for the AgriPredict project moving forward. Probably the most notable one would involve substituting the artificial training dataset with data collected from various real-world sources including national agro-databases, field IoT devices, and satellite image pipelines. Introduction of additional metrics such as NDVI values obtained from satellite sensors, e.g., Sentinel-2 or Landsat-8 would allow for taking into account spatial conditions of crop development that would be impossible to introduce otherwise.

When it comes to modeling, combining predictions generated by all three models: Random Forest, XGBoost, and Tuned Random Forest using a meta-learner could possibly help overcome the 0.9117 R^2 barrier. An option to consider would be a hybrid DL architecture that will combine convolutional and gradient-boosted layers in order to process both structured agronomic and unstructured satellite imagery datasets.

The Retrain Model tab present in the existing system design hints at a broader vision that involves creating a self-learning system capable of assimilating new information based on observations recorded in fields by farmers as new observations are provided to the system. Such a system can be developed through federated learning, allowing updates to be made to the model without needing to collate all farmer data at a central server. The last addition that should be considered is a natural language explanation of each recommendation from the predictive model through a large language model.

XII. CONCLUSION

Retrain Model tab already present in the current interface highlights the fact that the system can be made to learn continually by including new observations from the field uploaded by farmers and adapting itself according to the evolving agricultural conditions in the region. In such an environment, federated learning can prove to be a useful approach as the update of the model would not require the

consolidation of sensitive information related to individual farms in one server location. A final improvement that could be considered is the generation of natural language explanation corresponding to each model output using a large language model. This paper has provided an overview of AgriPredict, which is a Crop Yield Intelligence System using supervised machine learning algorithms for predicting agricultural yield. Using three ensemble machine learning models like Random Forest, XGBoost, and tuning hyperparameters of Random Forest in a dataset with one million entries covering a variety of crops and environments, the system has been found to make predictions with very high accuracy. Among these models, XGBoost performed the best with an R^2 of 0.9117 and MAE of 0.402 t/ha..

The attribute importance analysis of the feature selection algorithm used in the tool indicates that rainfall accounts for 60.7% of yield results while fertilizer use follows at 20.1% and irrigation comes in third with 12.6%. These results are in line with current agricultural theory, thus further supporting the effectiveness of what the model learned from the data. The web application provided along with AgriPredict ensures that users without technical skills can benefit from machine learning.

AgriPredict provides a demonstration of the potential applications of machine learning as long as due consideration is given to quality of training data, choice of a proper model and the interpretability of output. The software is designed to be especially useful for countries that still have to develop their agriculture because of high costs of expert consulting and limited impact of any improvements to yield forecasting accuracy on overall food supply.

Given that the world's agriculture is facing increasing pressures due to changing climate patterns, increasing population, and limited natural resources, technology such as AgriPredict will emerge as a major component of the precision farming industry in the coming years. With continuous improvement, incorporation of real-world data, and implementation on an agricultural policy scale, technologies of this nature can truly help create a more food-sufficient and agriculturally resource-efficient world.

REFERENCES

- [1] K. G. Liakos, P. Busato, D. Moshou, S. Pearson, and D. Bochtis, "Machine learning in agriculture: A review," *Sensors*, vol. 18, no. 8, p. 2674, 2018.
- [2] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, San Francisco, CA, 2016, pp. 785–794.
- [3] J. H. Jeong, J. P. Resop, N. D. Mueller, D. H. Fleisher, K. Yun, E. E. Butler, et al., "Random forests for global and regional crop yield predictions," *PLOS ONE*, vol. 11, no. 6, e0156571, 2016.
- [4] P. Filippi, E. J. Jones, N. S. Wimalathunge, P. D. Somarathna, L. E. Pozza, S. A. Ugbaje, et al., "An approach to forecast grain crop yield using multi-layered, multi-farm data sets and machine learning," *Precision Agriculture*, vol. 20, no. 5, pp. 1015–1029, 2019.
- [5] A. Nevavuori, N. Narra, and T. Lipping, "Crop yield prediction with deep convolutional neural networks," *Computers and Electronics in Agriculture*, vol. 163, 104859, 2019.
- [6] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, et al., "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [7] FAO, "The future of food and agriculture: Trends and challenges," Food and Agriculture Organization of the United Nations, Rome, 2017.
- [8] S. Krithika and M. Prakash, "Prediction of crop yield using machine learning algorithms," *Int. Journal of Engineering and Advanced Technology (IJEAT)*, vol. 9, no. 3, 2020.