

SUB-ID-738031

By Kush .

WORD COUNT

5218

TIME SUBMITTED

24-APR-2026 05:34AM

PAPER ID

121300443

Crop Yield Intelligence System (AgriPredict)

27

A Machine Learning-Based Approach to Agricultural Yield Prediction

1

Kush
Department of Computer Science & Engineering
Lovely Professional University
Phagwara, India
sonisingh0107@gmail.com

Nishar

Department of Computer Science and Engineering
Lovely Professional University
Phagwara, India
sonisingh0107@gmail.com

April 2026

Abstract

Forecasting agricultural output has become central to tackling the growing challenge of global food security. With the world's population on a steady rise, there is an ever-increasing need for farming practices grounded in data and supported by modern analytical tools. This paper introduces AgriPredict, a Crop Yield Intelligence System built upon supervised machine learning to estimate crop productivity in tons per hectare using a broad range of environmental and agronomic inputs. The system incorporates three predictive models — Random Forest, XGBoost (Extreme Gradient Boosting), and a hyperparameter-optimized Random Forest variant — all trained on a large-scale synthetic dataset comprising one million agricultural records that cover four geographic regions, six soil categories, six crop varieties, and three weather scenarios.

2

Among the models evaluated, XGBoost delivered the strongest results, recording a coefficient of determination (R^2) of 0.9117 and a mean absolute error (MAE) of 0.402 tons per hectare. These figures indicate that the model accounts for more than 91% of the variation observed in crop yield outcomes. An analysis of feature contributions shows that precipitation is the dominant factor at 60.7%, with fertilizer application contributing 20.1% and irrigation accounting for 12.6%. The platform offers a browser-based interface through which users can enter field-specific conditions and instantly obtain yield forecasts, making it practical for farmers, planners, and agricultural researchers. This paper covers the system design, data preparation workflow, rationale for model selection, the training and evaluation approach, and the broader implications of real-world deployment.

Keywords: crop yield forecasting, XGBoost, Random Forest, machine learning, smart farming, food security, feature contribution, regression analysis.

I. INTRODUCTION

Farming serves as the economic foundation of much of the developing world. In countries such as India, more than 58% of rural communities depend on agriculture as their main source of income. Yet despite this critical role, the sector is persistently exposed to challenges such as unpredictable weather patterns, uneven soil conditions, and a lack of structured, evidence-based planning tools. In regions where smallholder farmers make up the bulk of the agricultural workforce, the absence of reliable yield forecasting mechanisms results in poor resource management, overuse of fertilizers and water, heightened financial uncertainty, and ultimately, threats to food availability at both regional and national levels.

Conventional methods of estimating crop yields — typically carried out through manual field inspections, expert agronomic judgment, or basic rule-of-thumb models — tend to be both time-consuming and prone to significant error. The rise of machine learning as a transformative force in smart farming has introduced entirely new possibilities for developing automated, scalable, and high-accuracy prediction systems. These algorithms are capable of processing vast quantities of multi-dimensional data and

uncovering intricate, non-linear relationships between input conditions and harvest outcomes — patterns that traditional statistical approaches are generally too rigid to detect.

This paper presents AgriPredict, a Crop Yield Intelligence System that draws on a combination of machine learning algorithms to forecast harvest output from nine input parameters: rainfall, temperature, soil type, geographic region, crop variety, prevailing weather, fertilizer application, irrigation status, and duration to harvest. The system was developed and validated on a dataset containing one million records. The core motivation behind AgriPredict is to equip smallholder farmers and agricultural policymakers with an easy-to-use, real-time planning tool that significantly cuts down on the time and cost of manual yield estimation. On a broader level, the system aims to support the United Nations Sustainable Development Goal 2 (Zero Hunger) by fostering data-guided farming decisions ahead of each growing season.

The rest of this paper is structured as follows: Section II surveys related work on machine learning applications in crop yield prediction. Section III outlines the overall system architecture. Sections IV and V describe the models used and explain how they function. Section VI presents the dataset and input features. Sections VII through X cover data

partitioning, the training process, performance results, and visual analysis. Sections XI and XII address the system's strengths, limitations, and future directions, while Section XIII provides concluding remarks.

II. LITERATURE REVIEW

Over the past decade, the use of machine learning in agricultural yield forecasting has grown substantially, fueled by advances in remote sensing technologies, the proliferation of IoT-based field sensors, and the increasing digitization of farming records. Researchers have tested a broad spectrum of approaches, from conventional regression-based techniques to complex deep learning architectures, with each method showing particular strengths depending on the characteristics of the data and the nature of the forecasting objective.

Liakos et al. (2018) carried out a wide-ranging review of machine learning in precision farming, cataloguing applications of support vector machines, neural networks, and decision tree-based methods in yield prediction tasks. A key takeaway from their study was that ensemble approaches, especially Random Forest, tended to outperform individual estimators by mitigating overfitting through bootstrap aggregation. In a related vein, Jeong et al. (2016) showed the effectiveness of Random Forest in estimating soybean and maize yields using NDVI values derived from satellite imagery across extensive agricultural zones in the United States.

XGBoost, first proposed by Chen and Guestrin (2016), has since established itself as a leading algorithm for structured, tabular prediction problems, both in competitive machine learning settings and applied scientific research. Its gradient boosting framework, which includes built-in regularization, generally achieves better generalization than a standard Random Forest on most regression tasks. Filippi et al. (2019) demonstrated this by applying XGBoost to forecast wheat yields from soil nutrient profiles and weather data, obtaining R^2 values above 0.88 — a finding that closely aligns with the results reported in this study.

More recently, deep learning architectures such as Long Short-Term Memory (LSTM) networks have been examined for sequential crop yield forecasting. Although these models are theoretically well-suited to capturing time-dependent patterns, they typically require considerably larger datasets and greater computational capacity. For structured, non-sequential datasets like the one used in this work, ensemble tree-based methods such as XGBoost and Random Forest continue to be the preferred choice in practice, offering a strong combination of predictive accuracy, interpretability, and computational efficiency.

Notwithstanding the volume of existing research in this domain, many developed systems are constrained by their narrow scope, having been built for particular crops or geographic regions, which restricts their broader applicability. AgriPredict seeks to overcome this limitation by building a multi-class, multi-region predictive model capable of generating yield estimates across a wide variety of

agricultural environments, making it a more versatile and broadly useful tool for precision agriculture.

III. SYSTEM ARCHITECTURE

AgriPredict is built as an end-to-end intelligent platform that integrates a data ingestion and preprocessing pipeline, a backend for model training and performance evaluation, and an interactive browser-based interface for user interaction. The system is organized around a three-tier architecture — consisting of a data layer, a model layer, and a presentation layer — to ensure a clear separation of concerns, ease of maintenance, and straightforward extensibility.

A. Data Layer

The foundation of the system is a structured CSV dataset containing one million agricultural entries. Each row captures a combination of environmental and agronomic conditions paired with the observed crop yield in tons per hectare. The data layer handles all upstream data operations, including loading the dataset from disk, resolving missing entries, encoding categorical variables into numeric form, scaling numerical features, and dividing the data into training and testing partitions.

B. Model Layer

This layer houses three regression models: a baseline Random Forest, an XGBoost Regressor, and a hyperparameter-optimized Random Forest. All models are implemented using the scikit-learn and XGBoost libraries in Python. Once trained, each model is serialized via the joblib module and saved to disk, enabling rapid loading during inference without the need to retrain. Evaluation is carried out on the withheld test partition, with metrics such as R^2 , RMSE, MAE, and cross-validation scores computed and recorded for comparison.

C. Presentation Layer

The frontend is a React.js web application running on a local development server at port 5000. It is organized into five tabs: Predict Yield, Analysis, Model Details, Retrain Model, and About Project. The prediction screen collects values for all nine input parameters and sends a request to the backend inference endpoint, displaying the returned yield estimate. The Analysis tab renders interactive bar charts built with Recharts, comparing yield outcomes across soil categories, geographic regions, and fertilizer or irrigation usage. The Model Details tab enables side-by-side comparison of model-specific metrics and learning curve visualizations.

D. Input–Processing–Output Pipeline

The complete workflow is initiated when the user submits farm-level parameters through the interface. These values are then validated, encoded, and normalized using the identical preprocessing pipeline employed during model training, ensuring consistency between training and inference. The resulting feature vector is passed to the chosen model, which generates a yield estimate. This predicted value is rendered on the screen and highlighted on the Actual vs. Predicted scatter chart to give the user a sense of where their

result falls within the test distribution. A persistent model badge in the navigation header shows the name of the currently active model along with its R^2 score and the total training dataset size.

4 IV. MACHINE LEARNING MODELS

A. Random Forest Regressor

Random Forest is a tree-based ensemble method that builds a large number of decision trees during training and combines their individual outputs by computing the average prediction for regression problems. Its underlying mechanism relies on bootstrap aggregation, or bagging, where each tree is built from a randomly drawn sample of the training data selected with replacement. Added source of variation comes from the fact that only a random subset of input features is evaluated at each branching point, which promotes diversity across trees and guards against overfitting. Within AgriPredict, the Random Forest model attained an R^2 of 0.8989, an RMSE of 0.5386, and a cross-validation mean R^2 of 0.8990, reflecting solid and consistent generalization behavior. Its learning curve showed stable performance across different training set sizes, with training and validation scores both settling near 0.90, which is characteristic of a well-regularized model with low variance.

B. XGBoost Regressor 22

XGBoost, short for Extreme Gradient Boosting, is a sequential ensemble method that trains one decision tree at a time, with each new tree designed to correct the errors left by its predecessors. In contrast to bagging-based techniques, boosting directs more attention toward samples that were predicted poorly, causing the model to progressively sharpen its focus on harder cases. XGBoost also embeds L1 and L2 regularization directly into its loss function to limit overfitting and includes native support for efficiently handling datasets with missing or sparse values. In AgriPredict, the XGBoost model was configured with 100 estimators and a maximum depth of 20 per tree. Trained on 800,000 records and evaluated on a separate set of 200,000, it achieved the top score across all performance metrics: $R^2 = 0.9117$ and MAE = 0.402 tons per hectare. Given this performance, XGBoost was designated as the primary inference model and is shown as the active model in the navigation header.

C. Tuned Random Forest

To determine whether structured hyperparameter tuning could narrow the accuracy gap between Random Forest and XGBoost, a refined version of the Random Forest was also built. The tuning process involved a cross-validated grid search spanning parameters such as the number of trees, maximum tree depth, minimum number of samples at a split, and the number of features sampled at each node. The optimized variant did show a slight gain over the untuned baseline, but it was still unable to match XGBoost, whose sequential learning framework is inherently more expressive for this type of problem. This head-to-head comparison establishes a clear empirical basis for model selection and

confirms XGBoost as the most appropriate choice for production deployment.


V. WORKING OF THE MODEL

A. Preprocessing Pipeline

Before being passed into any model, the raw agricultural data moves through a structured multi-step preprocessing routine. Categorical fields — specifically Crop Type, Region, Soil Type, and Weather Condition — are converted to numeric form through label encoding, assigning each distinct category a unique integer identifier. Numeric features, including Rainfall in millimeters, Temperature in degrees Celsius, Days to Harvest, Fertilizer Used (a binary flag), and Irrigation Used (also binary), are rescaled using min-max normalization to bring all values into the [0, 1] interval. This prevents features with inherently larger numerical ranges from exerting undue influence over the model's learning process.

B. Feature Engineering and Selection

Feature contribution scores were extracted using XGBoost's native attribution mechanism to understand which inputs drive predictions most strongly. The results showed that Rainfall is the most influential factor, responsible for 60.7% of the model's predictive weight, firmly establishing it as the primary determinant of crop productivity in this context. Fertilizer application ranks second at 20.1%, reflecting the well-documented agronomic reality that nutrient supplementation substantially boosts plant output. Irrigation usage follows at 12.6%, underscoring the importance of managed water supply in offsetting variable precipitation levels. Temperature and Days to Harvest contribute 2.8% and 1.6% respectively, while the combined contribution of categorical variables — Crop Type, Soil Type, Region, and Weather Condition — amounts to 2.3%. These findings are consistent with foundational agricultural science and reinforce confidence in the model's interpretability.



C. Prediction Process

When a prediction request is made, the values entered through the web interface are routed through the preprocessing pipeline. The pipeline applies the same encoding mappings and normalization parameters that were

established during training, maintaining full consistency between training-time and inference-time transformations. The resulting feature vector is then forwarded to the saved XGBoost model, which processes it through its ensemble of 100 gradient-boosted trees to arrive at a final yield estimate. The model returns a single continuous numerical value representing the expected crop output in tons per hectare. This result is displayed in the interface alongside a scatter chart that plots the current prediction against the distribution of actual and predicted values from the test set, offering the user useful context for interpreting the output.

VI. DATASET AND PARAMETERS

A. Dataset Overview

The dataset on which AgriPredict was trained and evaluated comprises one million artificially generated agricultural records. After removing duplicate entries and those containing unrealistic values during the data quality review, 999,769 records were retained for further use. The data covers four geographic zones (North, East, South, and West), six soil categories (Chalky, Clay, Loam, Peaty, Sandy, and Silt), six crop varieties (Soybean, Wheat, Rice, Cotton, Maize, and Barley), and three weather classifications (Sunny, Rainy, and Cloudy). Recorded yield values span a range from 0 to 9.96 tons per hectare, with an overall dataset mean of 4.65 tons per hectare.

B. Input Features

Nine input variables define each record in the dataset: (1) Crop Type, which identifies the planted variety and is represented as a categorical feature; (2) Region, indicating the farm's geographic zone (North, East, South, or West), which indirectly encodes climate and terrain characteristics; (3) Soil Type, reflecting the textural class of the farmland's soil, which in turn affects its capacity to retain moisture and deliver nutrients; (4) Weather Condition, capturing the dominant seasonal pattern (Sunny, Rainy, or Cloudy); (5) Rainfall in millimeters, spanning a range from 100 mm under dry conditions to 1000 mm under heavy precipitation; (6) Temperature in degrees Celsius, covering a band from 15°C to 40°C; (7) Days to Harvest, indicating the full growth cycle length between 60 and 149 days; (8) Fertilizer Used, a binary flag (Yes or No) showing whether chemical fertilizer was applied; and (9) Irrigation Used, another binary flag denoting whether supplemental irrigation was in use.

C. Hyperparameters


The XGBoost Regressor was configured with the following settings: `n_estimators = 100` (total number of boosting iterations), `max_depth = 20` (depth limit per tree), `learning_rate = 0.1`, `subsample = 0.8`, `colsample_bytree = 0.8`, `reg_alpha = 0.1` (L1 penalty), and `reg_lambda = 1.0` (L2 penalty). The baseline Random Forest was run using scikit-learn's default configuration, with only `n_estimators` set to 100. The Tuned Random Forest was subjected to a grid search across `n_estimators` in {50, 100, 200}, `max_depth` in {10, 20, None}, and `min_samples_split` in {2, 5, 10}.

VII. DATA SPLITTING AND VALIDATION

The 999,769 cleaned records were divided into training and testing sets using an 80/20 ratio, which is a standard convention in machine learning workflows. This produced roughly 799,815 samples for model training and 199,954 samples for evaluation. The partitioning was carried out with random stratification applied to preserve the overall yield distribution in both subsets, ensuring neither partition was skewed toward any particular yield range.

Beyond the basic train-test split, five-fold cross-validation was also applied to both the Random Forest and XGBoost models. In this procedure, the training data is partitioned into five equally sized segments. At each iteration, the model is trained on four of these segments and validated on the fifth, cycling through all possible combinations and averaging the results. This yields a more reliable and lower-variance estimate of true generalization performance. XGBoost's average cross-validation R^2 of 0.9117 matched its held-out test score closely, confirming minimal overfitting and strong out-of-sample generalization. The Random Forest's cross-validation mean R^2 of 0.8990 was likewise consistent with its test R^2 of 0.8989.

A learning curve for the Random Forest model, accessible through the Model Details section of the application, visualizes R^2 performance against increasing proportions of the training data (ranging from 0% to 9% of the full set). Both the training score, which hovered near 0.98, and the validation score, which settled around 0.90, remained essentially flat across all training fractions examined. This stability indicates that the model reached its performance ceiling early in training — behavior that is typical of well-regularized ensemble models when trained on very large datasets.



VIII. TRAINING AND TESTING

Training for all three models was carried out using Python 3.10, with scikit-learn 1.2 and XGBoost 1.7 as the core machine learning frameworks. The entire training process was run on a standard CPU environment, with XGBoost's built-in multi-threaded tree construction enabling efficient parallel processing across the full one million-record dataset.

The training procedure followed a sequential pipeline: the CSV file was read into a pandas DataFrame, categorical columns were transformed using label encoding, numeric columns were normalized with min-max scaling, and the feature matrix was separated from the target vector. Each model was then fitted by calling its fit() method on the training partition. After training, the model objects were saved to disk using joblib.dump(), allowing them to be loaded quickly at inference time without repeating the training step.


Model performance was assessed on the held-out test partition using four standard regression metrics: the coefficient of determination (R^2), root mean squared error (RMSE), mean absolute error (MAE), and cross-validation mean R^2 . R^2 quantifies the share of variance in the target variable explained by the model, with values approaching 1.0 reflecting a stronger fit. RMSE places heavier emphasis on larger errors due to the squaring operation, while MAE provides a straightforward and easy-to-interpret measure of average prediction deviation expressed directly in tons per hectare.

XGBoost outperformed both Random Forest versions across every reported metric, recording $R^2 = 0.9117$ and MAE = 0.402 t/ha. The untuned Random Forest obtained $R^2 = 0.8989$, RMSE = 0.5386, and MAE = 0.4297 t/ha. The relatively small gap between models suggests that the dataset has a clear and learnable structure, and that the practical performance ceiling for tree-based regressors on this data sits somewhere in the R^2 range of 0.91 to 0.92. The Tuned Random Forest offered a small improvement over the baseline but still fell short of XGBoost, providing further justification for selecting XGBoost as the production-grade inference model.

IX. GRAPHS AND RESULTS ANALYSIS


A. Actual vs. Predicted Scatter Plot

The primary chart on the Predict Yield tab is a scatter plot that maps actual yield values against those predicted by the XGBoost model, drawn from a randomly sampled subset of test records. The x-axis reflects the ground-truth yield in tons per hectare, and the y-axis shows the corresponding model estimate. Each data point represents a single test instance. In a theoretically perfect model, all points would fall precisely along the 45-degree diagonal. In AgriPredict's scatter plot, the points cluster tightly around this diagonal, showing especially strong alignment for mid-range yields between 3 and 7 t/ha. A modest number of points deviate from the diagonal at both the low end (below 1 t/ha) and the high end (above 8 t/ha), which is a typical and acceptable feature of models applied to naturally variable agricultural data. The user's most recent prediction is marked in red on the chart, enabling direct comparison with the overall test distribution.



B. Fertilizer and Irrigation Impact

The bar chart under the Analysis tab illustrates average crop yield across four management scenarios: fields without fertilizer, fields with fertilizer applied, fields without irrigation, and irrigated fields. The differences are quite notable: fields where fertilizer was used produced an average of approximately 5.37 t/ha, compared to 3.89 t/ha for those without — representing a yield gain of around 38%. Irrigated fields averaged roughly 5.28 t/ha, versus 4.03 t/ha for non-irrigated fields, a difference of about 31%. These numbers lend quantitative support to the well-established agronomic understanding that proactive management of nutrients and water supply is among the most decisive factors in crop performance. The especially pronounced gap between fertilized and non-fertilized conditions mirrors the feature importance finding that fertilizer usage accounts for 20.1% of the model's predictive contribution.




C. Yield by Soil Type

This chart compares mean yield across six soil categories: Chalky, Clay, Loam, Peaty, Sandy, and Silt. The values fall within a narrow band, ranging from around 4.50 to 4.66 t/ha, with Chalky soil showing the highest mean yield at approximately 4.65 t/ha and Clay and Silt posting slightly lower figures near 4.63 t/ha. The compressed range across soil types implies that soil classification by itself does not

exert a major independent effect on yield in this dataset; rather, it interacts with precipitation levels, fertilizer availability, and irrigation practices to produce its combined impact. That said, the relative ordering of soil types can still serve as a useful practical reference for farmers making crop selection decisions based on the soil profile of their land.

D. Average Yield by Region

The region-based comparison chart shows that the East zone records the highest average yield at around 4.66 t/ha, followed by West at 4.65 t/ha, South at 4.63 t/ha, and North at 4.61 t/ha. Although these differences appear small in absolute terms, they carry meaningful weight in the context of large-scale agricultural strategy and resource planning. The relative advantage of the East region likely stems from more conducive rainfall distributions and temperature ranges encoded within the dataset. This visualization is especially relevant for government agricultural bodies looking to pinpoint high-productivity zones for crop diversification efforts or targeted infrastructure investments in irrigation.



E. Model Performance Summary

The Model Details tab lays out a comparative view of all three trained models. XGBoost's edge is evident from its R^2 of 0.9117 versus the Random Forest's 0.8989. Although the margin of roughly 1.3 percentage points may seem negligible at first glance, it translates to hundreds of thousands of more accurately estimated yield outcomes when the system is applied at a national deployment scale. The Random Forest's learning curve shows steady generalization throughout, with no evidence of overfitting, while its $n_estimators$ validation curve can help identify the point beyond which adding more trees produces diminishing performance gains.

X. ADVANTAGES AND LIMITATIONS

A. Advantages

AgriPredict holds a number of meaningful advantages over conventional approaches to crop yield estimation. To begin with, the system generates predictions instantly — something that field-based assessment methods would typically take weeks to accomplish. Additionally, having been trained on one million records drawn from a wide range of geographic zones, climatic conditions, and farming practices, the model is substantially more representative and

broadly applicable than systems built for a single crop or specific region. The feature contribution framework further adds value by offering interpretable insights into which inputs matter most, helping both farmers and policymakers understand not just the expected yield but the factors driving it. The web-based interface is designed to be intuitive and requires no technical background — a farmer can enter basic field parameters and receive useful output immediately. Finally, the availability of three distinct models within the same platform gives users the ability to compare results side by side and choose the most appropriate algorithm for their specific deployment context.

At a broader social level, AgriPredict is well-positioned to support government agencies tasked with food security by providing early national-level crop output projections. It could also help financial institutions and insurers better quantify agricultural risk when structuring farm loans or crop insurance products. Non-governmental organizations working with low-yield farming communities could use the system to pinpoint the specific inputs most constraining productivity in those areas. These applications connect directly to the objectives of SDG 2 (Zero Hunger), as they enable evidence-based agricultural planning prior to the commencement of each growing season.

B. Limitations

Notwithstanding its strong benchmark results, AgriPredict is subject to a number of important limitations. Most notably, the underlying dataset is synthetically generated rather than collected from real agricultural environments or field sensors. Although synthetic data supports large-scale, controlled experimentation, it is unlikely to fully reflect the complexity of real farming conditions, including crop diseases, flood events, progressive soil degradation, and the cumulative effects of long-term climate change. Furthermore, the system currently handles only nine input variables, whereas practical precision agriculture applications routinely draw on dozens of additional factors such as soil pH levels, organic matter content, pest pressure measurements, and remote sensing-derived indices. A third limitation is the absence of confidence intervals or prediction uncertainty estimates — a feature that is particularly valuable in risk-sensitive planning contexts where decision-makers need to understand the reliability of forecasts alongside the point estimates.

XI. FUTURE SCOPE

A number of meaningful improvements are envisioned for future versions of AgriPredict. The most consequential of these would be the replacement of synthetic training data with records drawn from real-world sources such as national agricultural databases, field-deployed sensors, and satellite imagery pipelines. Integrating vegetation health indices like the Normalized Difference Vegetation Index (NDVI) from satellite platforms such as Sentinel-2 or Landsat-8 would enable the model to incorporate spatial crop condition information that purely tabular approaches cannot access.

On the modeling side, stacking the outputs of all three trained models — Random Forest, XGBoost, and the Tuned Random Forest — through a meta-learner could provide a path toward surpassing the current R^2 ceiling of 0.9117. Another direction worth exploring is a hybrid deep learning architecture that combines convolutional layers for extracting spatial patterns from satellite imagery with gradient-boosted layers for processing structured agronomic data, potentially unlocking accuracy improvements beyond what tree-based methods alone can achieve.

The Retrain Model tab already visible in the current interface points toward a longer-term goal: building a continuously learning system that can incorporate fresh field observations uploaded by farmers, allowing it to adapt over time as regional farming conditions evolve. Federated learning could support this vision by enabling model updates without requiring sensitive farm-level data to be centralized on a single server. A final enhancement worth prioritizing would be the integration of natural language explanations for each prediction, generated by a large language model, which would translate technical model outputs into plain-language recommendations that farmers with limited data experience can readily understand and act upon.

XII. CONCLUSION

This paper has described AgriPredict, a Crop Yield Intelligence System that leverages supervised machine learning to forecast agricultural output. Through the training and systematic comparison of three ensemble-based models — Random Forest, XGBoost, and a hyperparameter-tuned Random Forest — on a dataset of one million records encompassing a broad range of crop types and environmental settings, the system demonstrates high predictive accuracy. XGBoost emerged as the top performer, achieving $R^2 = 0.9117$ and $MAE = 0.402$ t/ha.

The feature attribution analysis embedded in the system identifies rainfall as the dominant influence on yield outcomes at 60.7%, with fertilizer application second at 20.1% and irrigation third at 12.6%. These findings are in line with accepted agronomic knowledge, which lends additional credibility to the patterns the model has internalized from data. The accompanying web application packages these capabilities in an approachable interface, making data-driven yield estimation available to users who lack a technical background.

AgriPredict illustrates how machine learning, when applied thoughtfully with care given to data integrity, sound model selection, and interpretable outputs, can meaningfully augment traditional agronomic expertise. The system is particularly well-suited for deployment in developing agricultural economies where expert advisory services are expensive or scarce, and where even incremental improvements in yield forecast accuracy can have substantial downstream effects on food availability and the financial wellbeing of farming households.

With global agriculture under growing strain from shifting climatic conditions, expanding populations, and

shrinking natural resources, intelligent predictive tools like AgriPredict are poised to become a core part of the modern precision farming toolkit. Through ongoing development, integration with real field data, and adoption at the level of agricultural policy, yield intelligence systems of this kind hold genuine promise for contributing to a world that is both better fed and more efficient in its use of finite agricultural resources.

REFERENCES

- [1] K. G. Liakos, P. Busato, D. Moshou, S. Pearson, and D. Bochtis, "Machine learning in agriculture: A review," *Sensors*, vol. 18, no. 8, p. 2674, 2018.
- [2] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining*, San Francisco, CA, 2016, pp. 785–794.
- [3] J. H. Jeong, J. P. Resop, N. D. Mueller, D. H. Fleisher, K. Yun, E. E. Butler, et al., "Random forests for global and regional crop yield predictions," *PLOS ONE*, vol. 11, no. 6, e0156571, 2016.
- [4] P. Filippi, E. J. Jones, N. S. Wimalathunge, P. D. Somarathna, L. E. Pozza, S. A. Ugbaje, et al., "An approach to forecast grain crop yield using multi-layered, multi-farm data sets and machine learning," *Precision Agriculture*, vol. 20, no. 5, pp. 1015–1029, 2019.
- [5] A. Nevavuori, N. Narra, and T. Lipping, "Crop yield prediction with deep convolutional neural networks," *Computers and Electronics in Agriculture*, vol. 163, 104859, 2019.
- [6] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, et al., "Scikit-learn: Machine learning in Python," *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [7] FAO, "The future of food and agriculture: Trends and challenges," Food and Agriculture Organization of the United Nations, Rome, 2017.
- [8] S. Krithika and M. Prakash, "Prediction of crop yield using machine learning algorithms," *Int. Journal of Engineering and Advanced Technology (IJEAT)*, vol. 9, no. 3, 2020.

8%

SIMILARITY INDEX

PRIMARY SOURCES

1	assets-eu.researchsquare.com Internet	62 words — 1%
2	www.mdpi.com Internet	62 words — 1%
3	Natasa Kleanthous, Abir Hussain. "Machine Learning in Farm Animal Behavior using Python", CRC Press, 2025 Publications	20 words — < 1%
4	Ahmed Saif, Md Tarikul Islam, Md Aktaruzzaman. "ATMeQ: A Machine Learning-Based Framework for Amyotrophic Lateral Sclerosis Disease using RNA-seq Meta-Analysis", Springer Science and Business Media LLC, 2026 Crossref Posted Content	18 words — < 1%
5	www.ivdon.ru Internet	18 words — < 1%
6	Abhilasha Sharma, Vishwas Rathi, Anupam Biswas, Anil Singh, Omer Rana. "Multimodal Artificial Intelligence in Precision Agriculture - Practices, Challenges, and Applications", CRC Press, 2026 Publications	15 words — < 1%
7	K. Hemachandran, Manjeet Rege, Zita Zoltay Paprika, K. V. Rajesh Kumar, Shahid Mohammad Ganie. "Handbook of Artificial Intelligence and Wearables - Applications and Case Studies", CRC Press, 2024 Publications	13 words — < 1%
8	ijsrem.com Internet	

13 words — < 1 %

9 [gulftime.ae](#)
Internet

11 words — < 1 %

10 [theaspd.com](#)
Internet

11 words — < 1 %

11 [www.frontiersin.org](#)
Internet

11 words — < 1 %

12 Bhargav Appasani, Amitkumar Vidyakant Jha.
"Engineering Applications of AI for Demand
Forecasting", CRC Press, 2026
Publications

10 words — < 1 %

13 [methodist.edu.in](#)
Internet

10 words — < 1 %

14 Abdullah, Syed Muhammad Fahad. "Machine
Learning Applications for Evaporation Predictions
From Small Reservoirs: Potential Water Savings for Lower Rio
Grande Valley, Texas.", The University of Texas Rio Grande
Valley
ProQuest

9 words — < 1 %

15 Bayou, Wuhib A.. "Modeling Ecological Uplift
Across Maryland's Urbanizing Watersheds: A Multi-
Stressor Machine-Learning Framework.", Embry-Riddle
Aeronautical University
ProQuest

9 words — < 1 %

16 Harvinder Singh, Priyanka Kaushal, Sarabjeet Kaur.
"Advanced Computing and AI for Sustainable
Environment, Energy, and Smart Manufacturing - AI-Driven
Sustainability Across Environment, Energy, and Industry", CRC
Press, 2026
Publications

9 words — < 1 %

17 [afd-prod-waf.cegedim.cloud](#)
Internet

9 words — < 1 %

-
- 18 ideas.repec.org Internet 9 words — < 1 %
-
- 19 s3-ap-southeast-1.amazonaws.com Internet 9 words — < 1 %
-
- 20 www.ijcmas.com Internet 9 words — < 1 %
-
- 21 Dinesh Goyal, Payal Bansal, Ritam Dutta, Madhav Sharma. "Intelligent Systems Using Semiconductors for Robotics and IoT - First International Conference, ICISRI-2024, Jaipur, India, November 29-30, 2024 Proceedings", CRC Press, 2025 Publications 8 words — < 1 %
-
- 22 Gartenberg, Rebecca. "Predicting the Success of Stock Trading Strategies Using Machine Learning", The Cooper Union for the Advancement of Science and Art, 2024 ProQuest 8 words — < 1 %
-
- 23 preview-www.nature.com Internet 8 words — < 1 %
-
- 24 www.bio-conferences.org Internet 8 words — < 1 %
-
- 25 R. N. V. Jagan Mohan, B. H. V. S. Rama Krishnam Raju, V. Chandra Sekhar, T. V. K. P. Prasad. "Algorithms in Advanced Artificial Intelligence - Proceedings of International Conference on Algorithms in Advanced Artificial Intelligence (ICAAAI-2024)", CRC Press, 2025 Publications 7 words — < 1 %
-
- 26 Thomas van Klompenburg, Ayalew Kassahun, Cagatay Catal. "Crop yield prediction using machine learning: A systematic literature review", Computers and Electronics in Agriculture, 2020 Crossref 7 words — < 1 %

27

Bui Thanh Hung, M. Sekar, Ayhan Esi, R. Senthil Kumar. "Applications of Mathematics in Science and Technology - International Conference on Mathematical Applications in Science and Technology", CRC Press, 2025

Publications

6 words — < 1%

28

Gholamreza Eslamifar, Ilya Zaslavsky, Hamid Balali, Vytautas Samalavičius, Alexander Fernald. "Forecasting change in groundwater storage: A multilayer perceptron model leveraging system dynamics simulation data", Journal of Hydrology: Regional Studies, 2026

Crossref

6 words — < 1%

29

Uche Onyekpe, Vasile Palade, M. Arif Wani. "Recent Advances in Deep Learning Applications - New Techniques and Practical Examples", CRC Press, 2025

Publications

6 words — < 1%